


Accumulators and their applications to Mimblewimble

Benedikt Bünz

Joint work with Ben Fisch, Dan Boneh


Grincon, January 2019

Mimblewimble

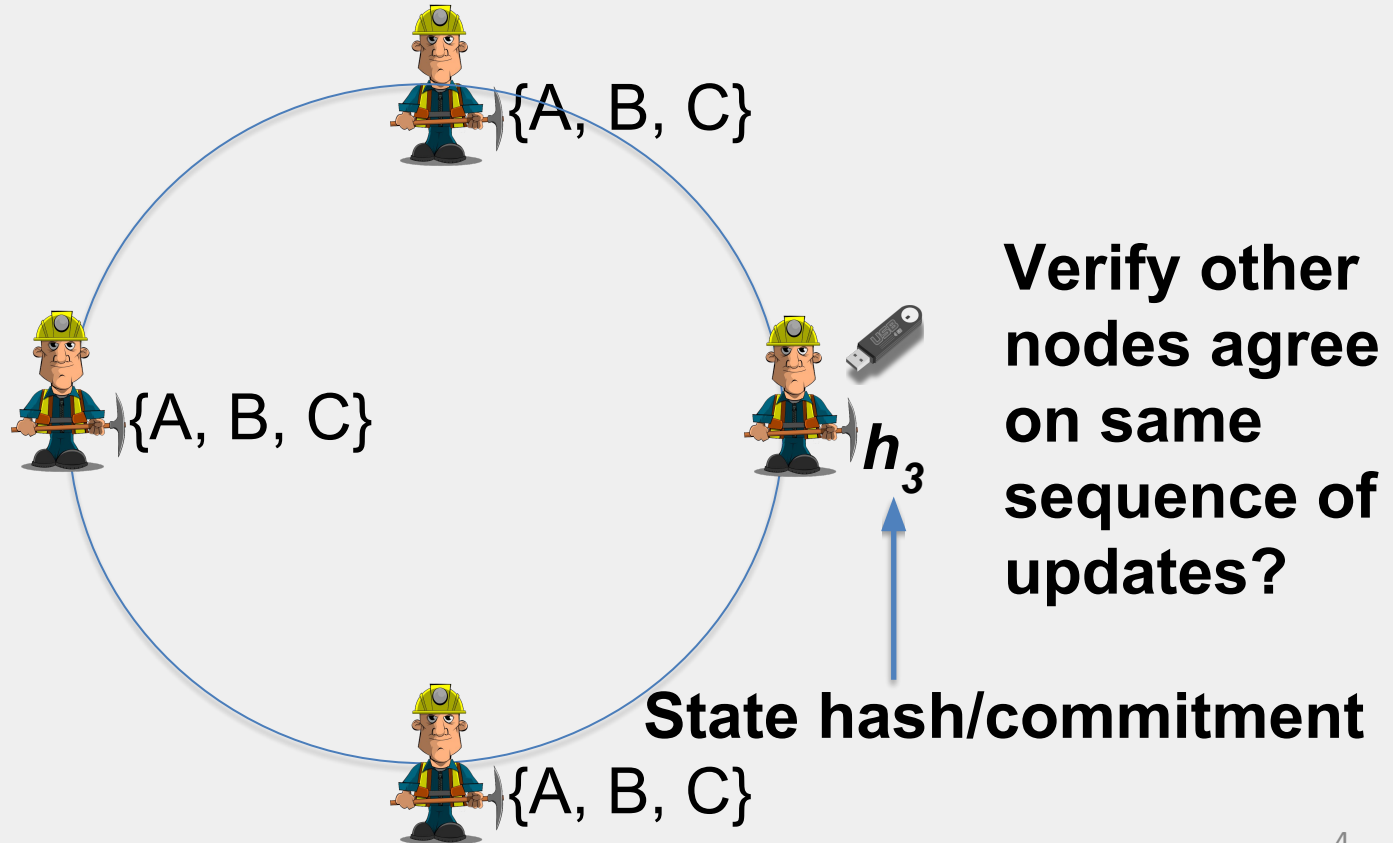


State is (almost) self verifying


Bitcoin State: A Growing Problem



Stateless verification




Stateless verification




**Stateless
node can't
participate in
consensus
with rules...**

Stateless consensus



Stateless consensus



How can
nodes check
validity of
update?


Consensus and State



Consensus

Checks proof and updates state


State Update + Proof



State storage

Provides proofs


Stateless consensus



How is the proof generated?
Does it require state?


Who generates the proof?

Stateless consensus



Bitcoin/Mimblewimble UTXOs

Unspent transaction outputs (Coins)



UTXOs

Miners agree on UTXO set S

$$h = \text{Commit}(S)$$

Accumulator

Transaction



Proof

$$h = \text{Commit}(S)$$
$$\text{UTXOs} \in S$$

Accumulators [Bd94, CL02]

Short
Membership
Witness

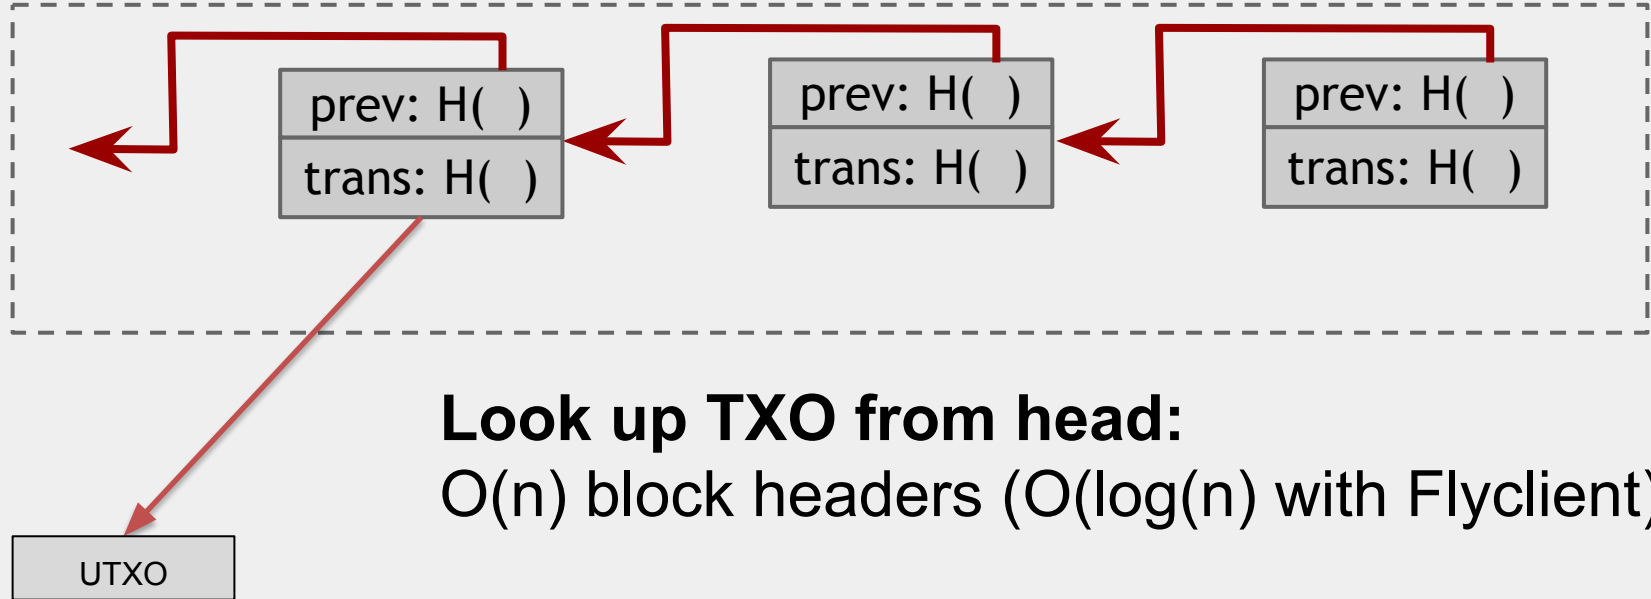
$A_{i+1} \leftarrow \mathbf{Add}(A_i, x)$
 $\pi = \mathbf{InclusionProof}(A_{i+1}, x)$
 $\mathbf{Verify}(\pi, A_{i+1}, x) = \{0,1\}$

Constant
size

Examples: Merkle trees
RSA Accumulators [CL02]
Pairing-based accumulators [NG05]

Bitcoin UTXOs

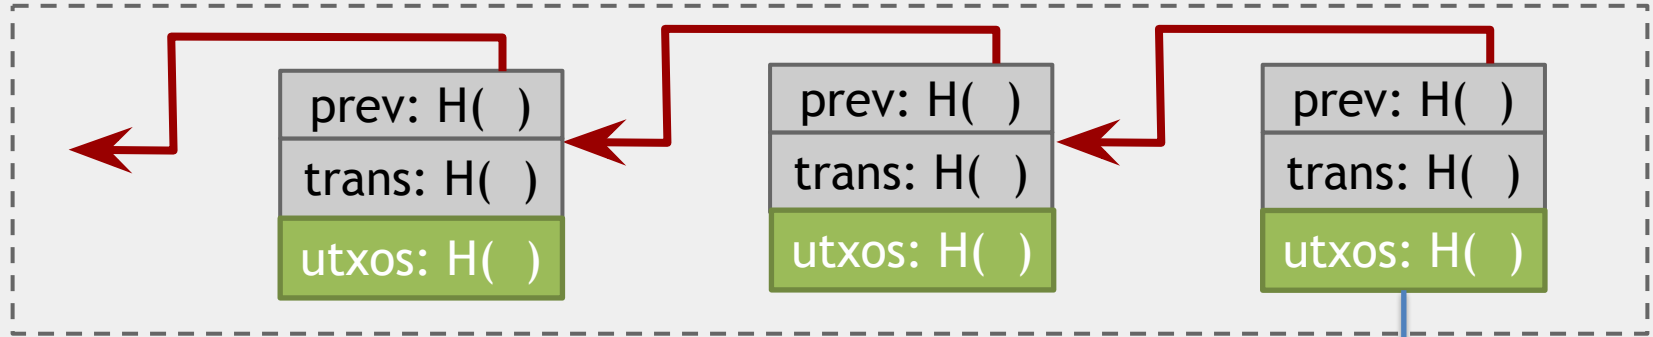
Unspent transaction outputs



Look up TXO from head:
 $O(n)$ block headers ($O(\log(n))$ with Flyclient)

Look up UTXO: All transactions

UTXO Commitments [TMA13]



Consensus ensures:
All UTXO committed here


Merkle Trees

Inclusion: $O(\log(n))$

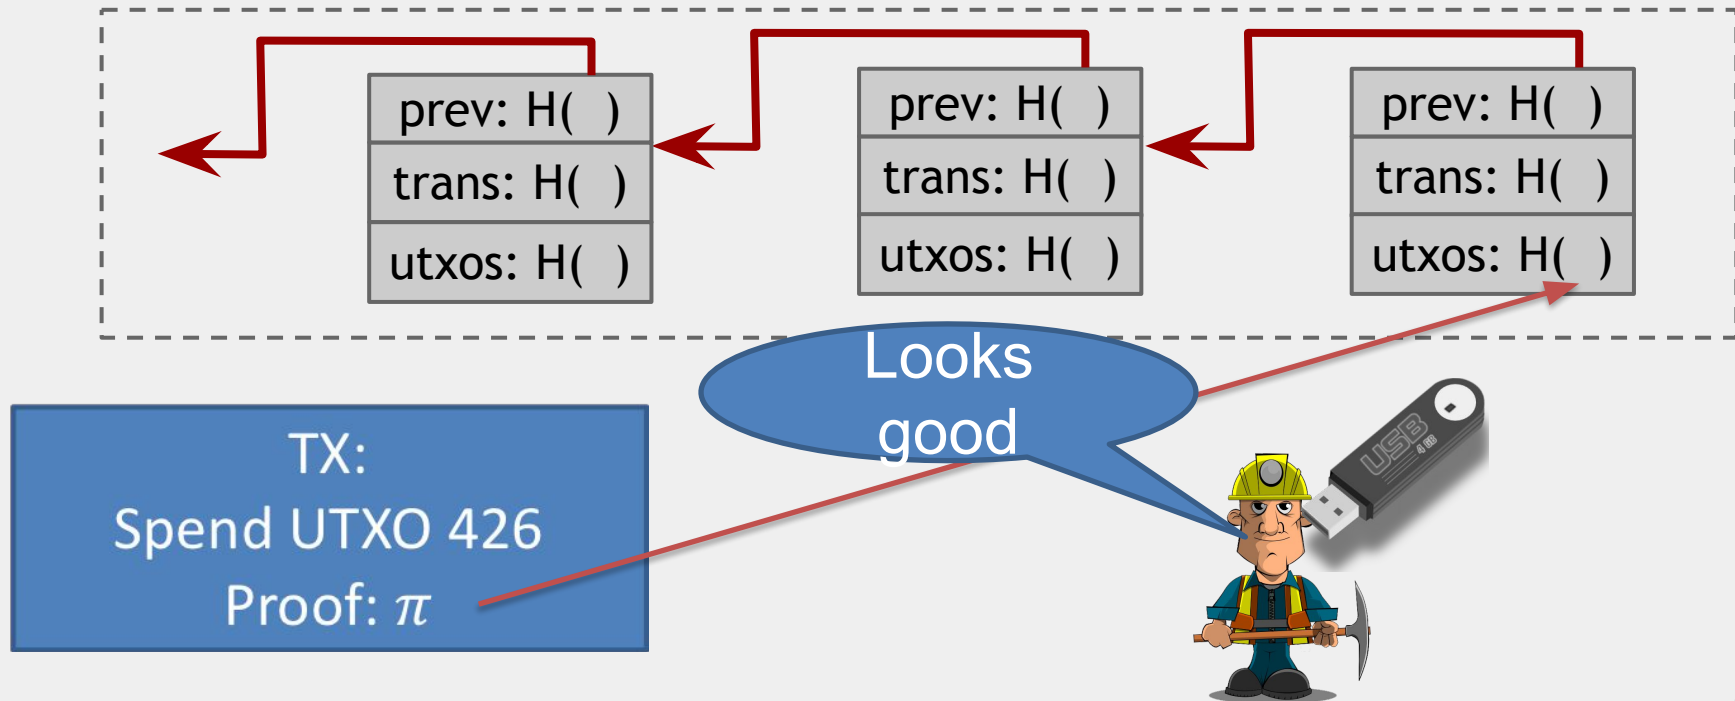
Exclusion: $O(\log(n))$ ¹

Update: $O(\log(n))$

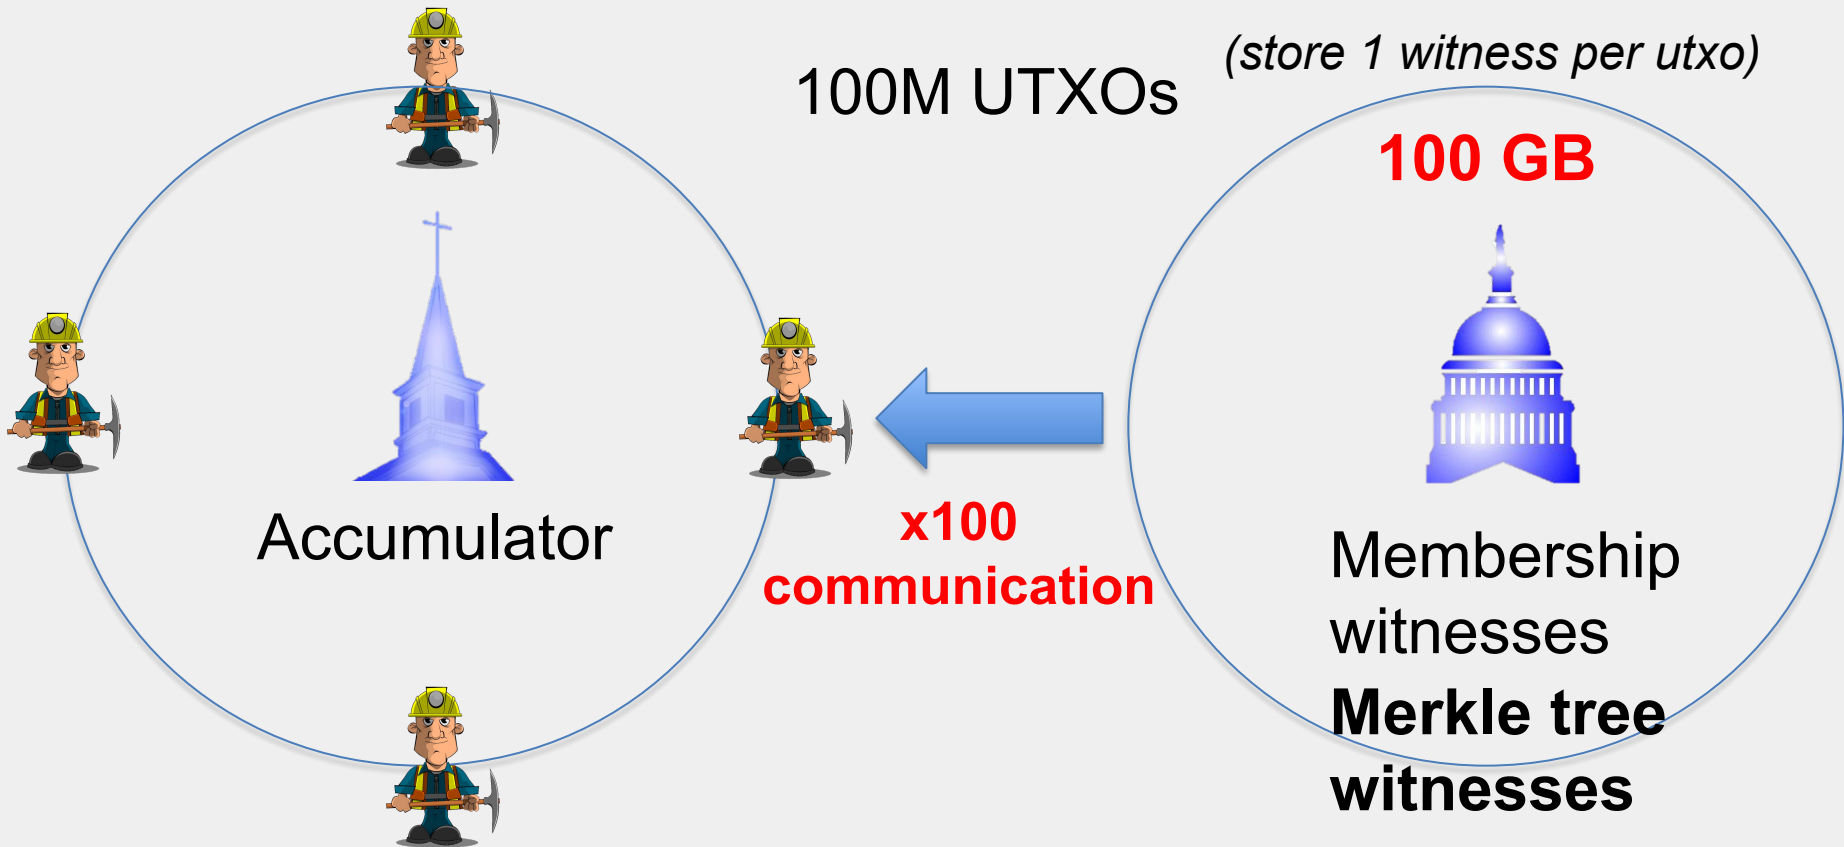
¹ If sorted



Stateless Full Nodes/Mining



Membership Witnesses



Membership Witnesses

Desiderata

- **Short** membership witness per item
- **Efficiently** updatable (storage + computation)
- **Efficient** verification
- **Aggregate witnesses?**
- **Batch generate & verify?**



Membership
witnesses

Problems with Merkle Trees

- $\text{Log}(n)$ inclusion proof per transaction
- Inclusion proofs can hardly be aggregated
 - 600 GB naively
 - 160 GB with many optimizations
- Verification not that cheap
 - Full node sync too slow
 - Proposed for only old transactions

RSA Accumulators [CL02, LiLiXue07]

Setup:

- Choose $N=pq$ where p, q are secret primes
- H : Hash function to primes in $[0, 2^\lambda]$
- $A_0 = g \in Z_N$ (*initial state*)

Add(A_i, x)

- $A_{i+1} = A_i^{H(x)}$

Del(A_i, x)

- $A_{i+1} = A_i^{1/H(x)}$

State after set S added:

$$u = \prod_{s \in S} s$$
$$A_t = g^u$$

Accumulator Proofs

InclusionProof(A,x):

- $\pi = A^{\frac{1}{x}} \in \mathbb{G}$
- Computed using trapdoor(p,q) Or $O(|S|)$

Verify(A, x, π)

- $\pi^x = A$

Efficient stateless updates to
(non)-membership witnesses:
[LiLiXue07]

Exclusion(A, x)

- $A = g^u$
- $a \cdot x + b \cdot u = \gcd(x, u) = 1$
- $\pi = (g^a, b) \Rightarrow \text{Verify } \pi^x \cdot A^b = g$

RSA Accumulator State of Art

Positives

- Constant size inclusion proofs (≈ 3000 bits)
Better than Merkle tree for set size > 4000
- Dynamic stateless adds (can add elements w/o knowing set)
- Decentralized storage (no need for full node storage)
 - *Users maintain their own UTXOs and membership proofs*

Room for improvement? **New work**

- Aggregate/batch inclusion proofs (many at cost of one)
- Trapdoor free (no trusted setup)
- Stateless deletes
- Faster (batch) verification

Batching Techniques for Accumulators with Applications to IOPs and Blockchains

Joint work with: Ben Fisch and Dan Boneh

<https://eprint.iacr.org/2018/1188>

Aggregate Membership Witnesses

$$\pi_1^x = A, \pi_2^y = A$$

Shamir's Trick:

$$a \cdot x + b \cdot y = 1$$

$$\pi_{1,2} = \pi_1^b \pi_2^a$$

$$\pi_{1,2}^{x \cdot y} = A$$

All membership witnesses per transaction block: ~3000 bits

RSA = Trusted Setup?

$N=p*q$, p,q unknown

Efficient delete needs
trapdoor

You can find N s in the wild
(Ron Rivest Assumption)

Class Groups [BW88,L12]

$CL(\Delta)$ – Class group of quadratic number field $\mathbb{Q}(\sqrt{\Delta})$
 $\Delta = -p$ (a large random prime)

Properties

- Element representation: integer pairs (a, b)

$$|a| \approx |b| \approx \sqrt{-\Delta}$$

No trusted
setup

- Tasks believed to be hard to compute:

Odd prime roots

Group order

- $\Delta \approx 1536 \text{ bits} \Rightarrow 128 \text{ bit security}$

Stateless Deletion

Delete with trapdoor(A_t, x):

- $A_{t+1} = A_t^{\frac{1}{x}}$

Using knowledge
of p, q

Delete with inclusion proof(A_t, x, π)

- $A_{t+1} = \pi$;

$$\pi = g^{\frac{u}{x}}$$

BatchDelete(A_t, x, y, π_1, π_2)

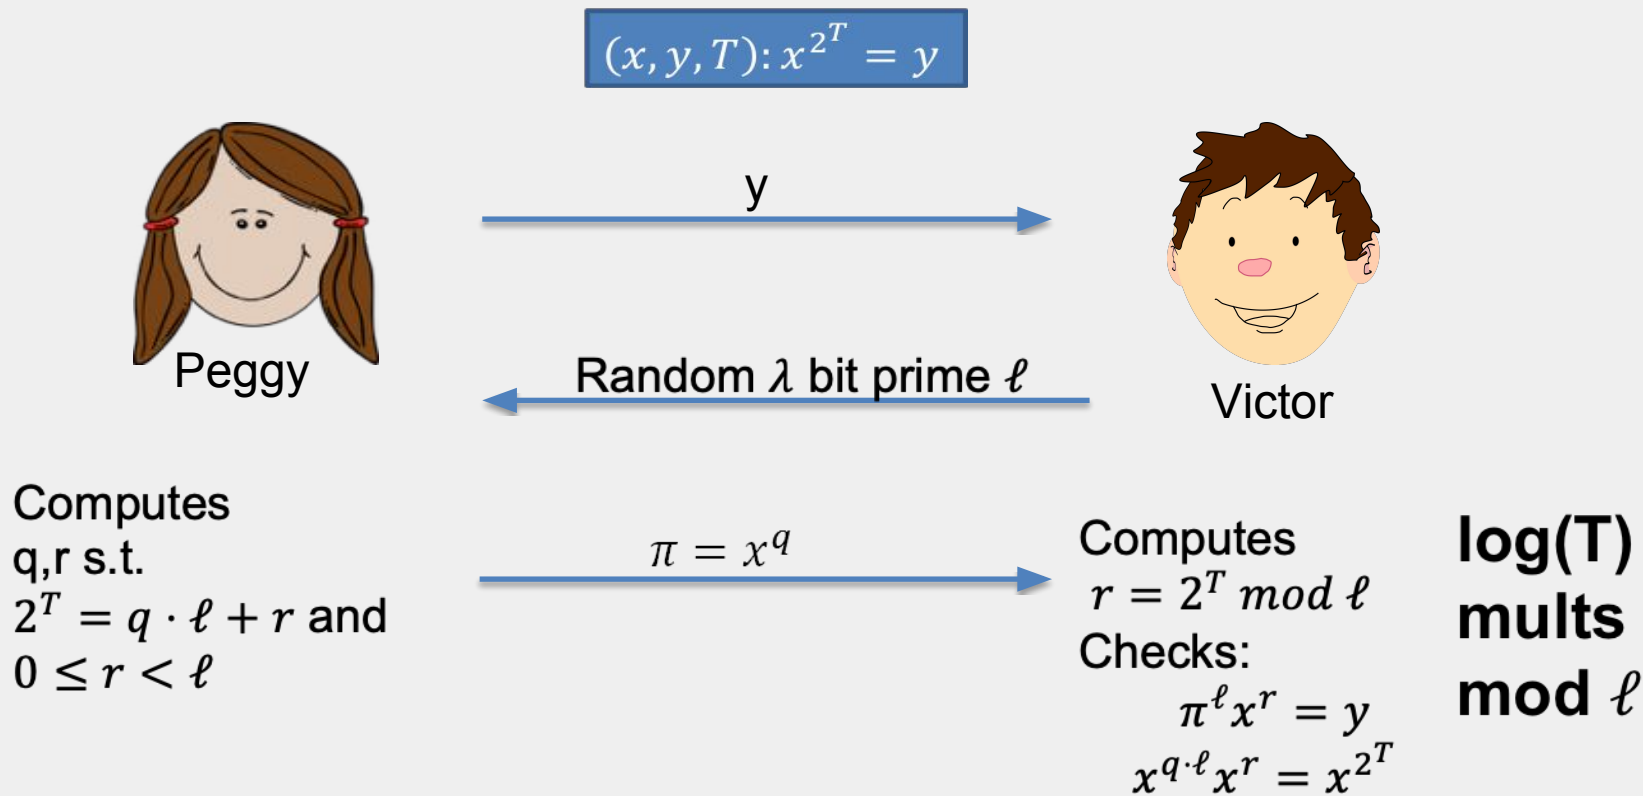
- Compute $\pi_{1,2}$ s.t. $\pi_{1,2}^{x \cdot y} = A_t$
- $A_{t+1} = \pi_{1,2}$

No State,
no Trapdoor,
asynchronous

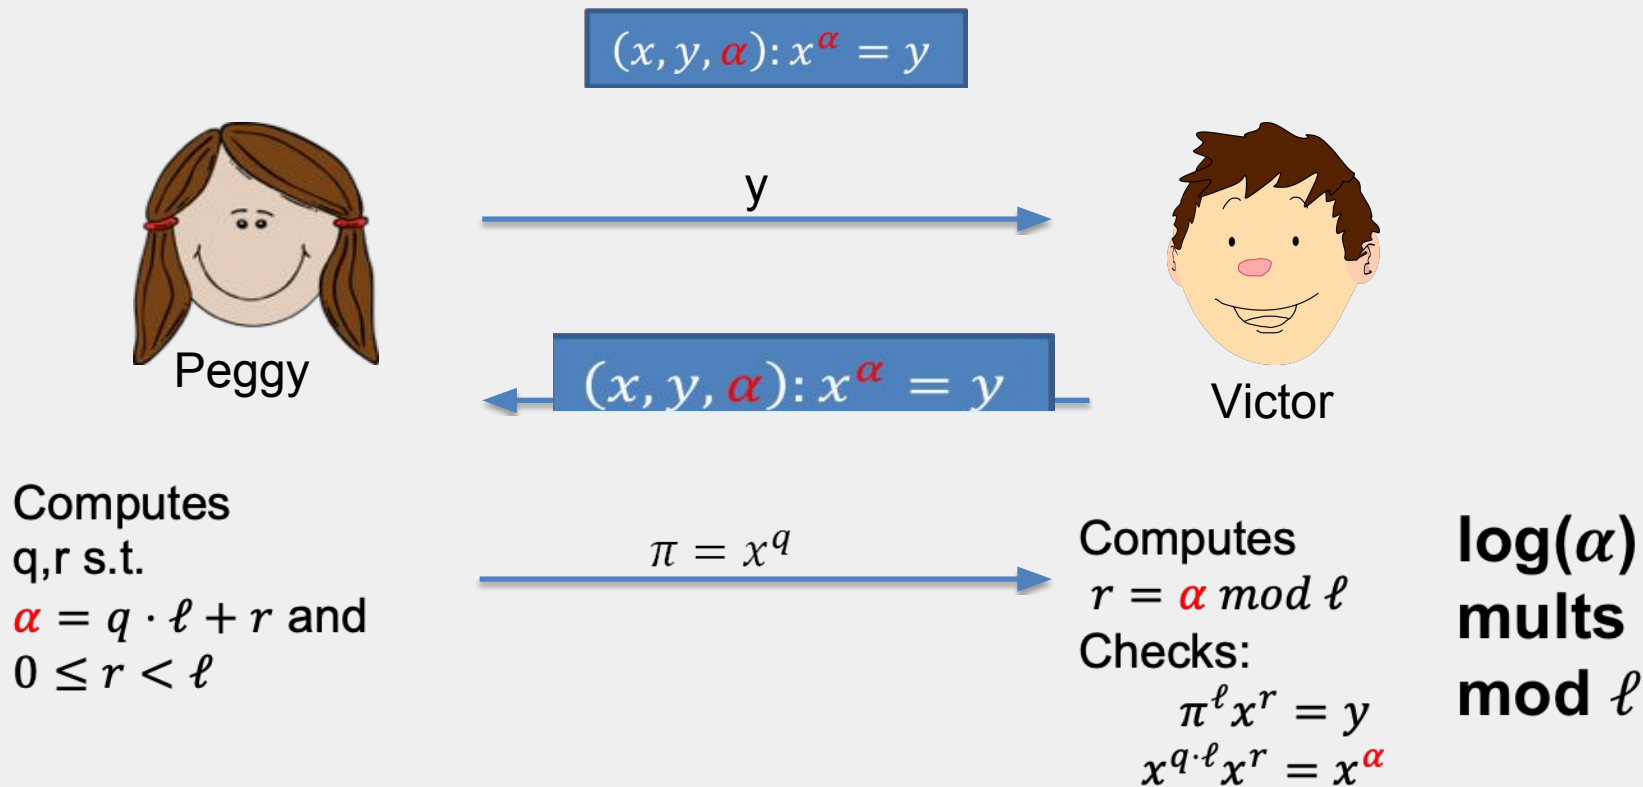
Verification of Witnesses Too slow?

- Java Big Integer Microbenchmark:
 - 600 exponentiations per second (256 bit exponents)
 - Verification/Full sync would be problematic*(On my laptop)*
- Class groups: No good benchmarks yet

Wesolowski Proof [Wesolowski'18]



Proof of Exponentiation (PoE)



PoE Efficiency

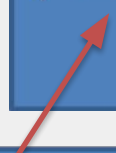
Both linear in
bitlength

$$(x, y, \alpha): x^{\alpha} = y$$

Direct
Verification:
 $x^{\alpha} = y \in \mathbb{G}$



PoE Verify:
 $r = \alpha \bmod l$
 $\pi^l g^r$



**Much
faster**

Exponentiation in \mathbb{G} vs. 128 bit long-division:
5000x difference for 128 bit security

PoE Assumption

Adaptive Root assumption:

For all efficient adversaries A

Given $u \in G \xleftarrow{\$} A(\lambda), \ell \xleftarrow{\$} \text{Primes}(\lambda)$

A will produce $w \in G$

$$w^\ell = u$$

With only negligible probability

G can't have any element of known order other than 1

-> We need to work over $Z_N^+ := Z_N / \{\pm 1\}$

Classgroups seemingly satisfies this property

Fast Block Verification



Header:
TXs: Spent S, new N
Signatures σ
A'_t, A_{t+1}, PoE



A_t

BatchDelete

A'_t

A'_t

BatchAdd

A_{t+1}

Verify σ
Verify PoE for BatchDel
Verify PoE for BatchAdd

Fast Full Sync verification



Mimblewimble:

UTXOs

Signatures σ

Range proofs π

A



Batch verify σ
Batch verify π
Verify PoE for BatchAdd

Performance

Macbook, Java BigInteger, JDK Hash

Merkle Tree: 26 x SHA-256:

$8.5 \mu\text{s} > 100,000/\text{s}$

Add: $g^x \bmod N$, $|x|=256$ bit $|N|=3072$:

$1535 \mu\text{s} > 600/\text{s}$

Verify: $x \bmod I$, $|x|=256$ bit $|I|=128$ bit

$0.3 \mu\text{s} + 50 \mu\text{s}$ for primality checking

$\sim 20,000/\text{s}$



Classgroups?

Takeaway Points



Consensus

**Shifting work from
miners to users**

**Distribute the storage
(load balanced
blockchain)**



State storage

References

- CL02: Camenisch Lysanskaya 2002 Dynamic Accumulators
- LiLiXue07: Li, Li, Xue 2007 Universal Accumulators
- CF: Catalone Fiore: Vector Commitments
- Todd:
<https://petertodd.org/2016/delayed-txo-commitments#further-work>
- MMR:
<https://github.com/opentimestamps/opentimestamps-server/blob/master/doc/merkle-mountain-range.md>
- UTXO: <https://bitcointalk.org/index.php?topic=101734.0>
- BW88: Buchmann and Williams